

# HDFS Container

For testing purposes, we have made a HDFS docker image that can be run in your local machine following this steps:

```
#You need to login in the docker registry before. If you don't, ask QA team.
docker pull qa.datio.mx:5000/datiomx/hdfs:0.0.1
docker run -dit --name hdfs -p 9000:9000 -p 50010:50010 -v /tmp/hdfs:/home/root/data/hdfs --hostname hdfs.dev.datio.com qa.datio.mx:5000/datiomx/hdfs:0.0.1
```

Now, check that you can make any operation to HDFS. First, add the following line to your `/etc/hosts` file:

```
127.0.0.1 hdfs.dev.datio.com
```

Then, configure your `core-site.xml` with the following configuration:

## core-site.xml

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://hdfs.dev.datio.com:9000</value>
  </property>
</configuration>
```

Now you can use the `hdfs` utility to check that all works fine:

```
./hdfs dfs -mkdir /user
./hdfs dfs -copyFromLocal <any_file> /user
./hdfs dfs -cat /user/<any_file>
```

Then you can destroy the container and run again, checking that your data is still there:

```
docker rm -f hdfs  
./hdfs dfs -cat /user/<any_file>
```

If you want to get an empty HDFS again, you need to delete /tmp/hdfs directory.